

Using Gated Recurrent Unit (GRU) Neural Networks to predict Google Stock prices

Julian Cabezas Pena Student ID: a1785086
Assignment 3. Deep Learning Fundamentals, 2020
University of Adelaide
`julian.cabezaspena@student.adelaide.edu.au`

Abstract

1. Introduction

In the

The prediction of stock prices in the future is a topic that has attracted both academic and industry investigators [11], and is often considered as one of the most challenging problems in the field of time series prediction [5]. Given the high volatility, complexity and nonlinearity of the stock market, trading can not rely on the trader's intuition or experience, making the use of machine learning methods a research hotspot [7]

In the last years, one of the methods that has been more frequently used to predict stock prices are the Recurrent Neural Network (RNN). These kind of neural network architecture deal with the problem of varying dimensionality, preprocessing sequences of input data of varying length, and are commonly used in language and audio processing [9], while also being suitable for time series prediction [11]. While other architectures such as the convolution neural networks or the multi layer perceptron present connections that push the information forward, Recurrent Neural Networks contain connections that bring the information backward to a layer as an input [9]

The most commonly used Recurrent Neural Network architecture is the Long Short-Term Memory network, that was introduced in 1997 by Hochreiter and Schmidhuber [4], that solved the problem of vanishing or exploding gradients at that time, introducing a backpropagation method that enforced a constant error flow through the units of the neural network [4]. The LSTM adds or removes information of the networks through the use of gates, that *remember* or *forget* information about the previous state in the recurrent network, a typical LSTM network includes a forget gate, an input gate, an output gate and memory cell [3].

Most recently, Cho *et al* [2] introduced a slightly different

version of Recurrent Neural Network, called Gated Recurrent Unit (GRU), that similarly to the LSTM, presents gates that control the flow of information, but not presenting a separate memory cell. Thus, the LSTM presents a controlled exposure of the of the memory content, while the GRU exposes all the content of the previous state [3].

Studies conducted by Chung *et al* [3] report similar performance between LSTM and GRU in various sequence modelling task in polyphonic music and speech signal dataset. Similarly, Zhao *et al* [11] report similar results in the field of stock price trend prediction between these two kind of architectures. One advantage of the GRU over the LSTM is that it contains a smaller number of parameters[3], usually making the training time slightly shorter.

The objective of this research is to implement a GRU architecture to predict the prices of the Google stock in the short future, implementing models to predict the next day Open, Low, High and Close price of the stock, along with the Volume of stocks being traded. Moreover, a model to predict the Low, High and Close price knowing the Open price of the same day will be tested to increase the usability of these prediction models in trading purposes.

2. Background

As abovementioned, the use of RNN in stock price prediction has been extensively reported in the scientific literature. Many of these papers use three forms of RNN: the *vanilla* RNN, the LSTM and the GRU for stock prediction, one clear example of this is the research conducted by Zhao *et al* [11], where these three architectures were tested to predict 180 stocks prices and indicators from the Shanghai Stock Exchange, the authors determined that the LSTM and GRU outperform vanilla implementations of RNN, also, they propose the inclusion of an attention mechanism, that increases the performance of the models.

Following a similar approach Qiu *et al* also propose the inclusion of an attention mechanism combined with a LSTM model to predict stock prices in several stock datasets or

indexes (S&P 500, DJIA and HSI). In this case, they used a wavelet transform to denoise the stock price data, along with data normalization. The paper aims to find the Open price of the stocks, also confirming that the inclusion of an attention layer can benefit the accuracy of the prediction.

Other approaches have also been attempted to predict the stock prices using different machine learning tools. As an example, Kara *et al* tested the use of a three layered feed-forward artificial neural network and a Support Vector Machine Model to predict the stock price of a sample of the Istanbul Stock Exchange, in this case, the authors aim to predict the price movement rather than the price itself.

Even though it is possible to say that most modern attempts to predict stock prices are based on machine learning or deep learning techniques, the stock prices can be analyzed using more traditional time series analysis methods. As an example, Adebisi *et al* [1] used an ARIMA model to predict the prices of the New York and the Nigeria Stock Exchanges

3. Methods

3.1. Google Stock Price Dataset

The Google Stock Price dataset (<https://www.kaggle.com/rahulsah06/google-stock-price>) consists in two separate databases (train and test), that together contain 1578 samples of data, comprising the years between 2012 and 2016 (included) and a small fraction of the year 2017. The database contains the different stock prices of the GOOG stock along the days where the stock exchange is open (generally from Monday to Friday). The data contains 6 fields:

- Date: Date of the record, in MM/DD/YYYY format
- Open: Price of the stock at the beginning of the day
- High: Higher price recorded by the stock in the day
- Low: Lower price recorded by the stock in the day
- Close: Price of the stock at the end of the day
- Volume: Amount of stocks that were traded in that day

The total dataset is divided into the training data, that contains 1578 samples (between the 3rd of January 2012 and the 30th of December 2016) and the test data, that contains 20 records (between the 3rd of January 2017 and the 31st of January 2017)

3.2. Preprocessing of the data

In Figure 1, it is possible to notice that the Close Price of the GOOGL stock is significantly larger than the High Price, identifying an inconsistency in the data, that is not

present in other dataset where this stock is present (such as the DJIA 30 dataset: <https://www.kaggle.com/szrlee/stock-time-series-20050101-to-20171231>). Looking at the specialized press, the GOOGL stock was splitted in the 3rd April 2014, when this Close Price starts being consistent again. As the split considered giving 2002 stocks for every 1000 shared owned by the investor, the Close price was divided by 2.002 if it was larger than the High price. In the case that the Close price remained inconsistent (smaller than the Low price or bigger than the High price), it was replaced by the Low or High Price.

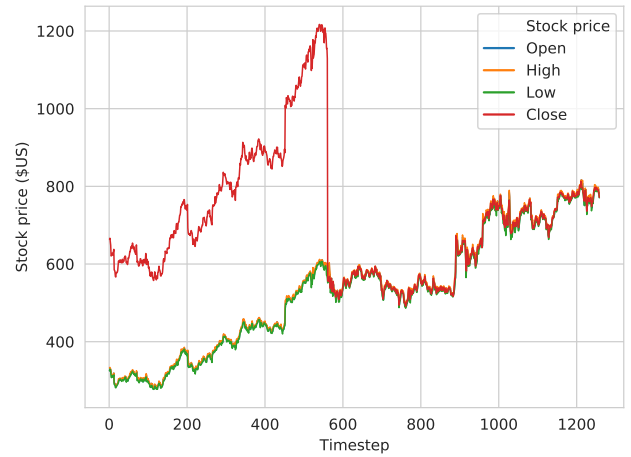


Figure 1. Open, High, Low and Close prices of the Google stock (uncorrected)

As seen in other implementations of deep learning for stock price prediction [11, 5], the stock prices were rescaled using minimum maximum normalization to leave the stock prices from 0 to 1, using the following equation:

$$x' = (b - a) \times \frac{x - \min(x)}{\max(x) - \min(x)} + a \quad (1)$$

where b is the new maximum (in this case +1), a is the new minimum (0), x is the original feature and x' is the transformed feature.

3.2.1 Sliding window

The prediction of time series data is usually performed using a sliding window or "look back" representation of the data. This means that the data that is used as input to predict the near future comes from the closest n days (usually around 30), as the data far away from the target usually does not provide useful information for the prediction [8].

To address this issue, a sliding window approach was implemented, generating sets of explanatory variables comprising the N days before the target data, and matching them with the target variable day, including the Open, Low,

High and Low prices, along with the trading Volume. For practical purposes, the Date field was neglected, and only the days where there were available records were considered as part of the window.

3.3. Gated Recurring Unit Neural Network

To predict the stock prices, a neural network comprised of Gated recurring neural networks (GRU) was used. The GRU is comprised by two gates (reset and update). On one side, the reset gate (r_t) is calculated as follows [3]

$$r_t = \sigma(W_r x_t + U_r h_{t-1} + b_r) \quad (2)$$

where x_t is the input data for time t , and h_{t+1} is the previous hidden state W_r and U_r are set of weights of the gate, and b_r is the bias term

While the update gate (z_t) is calculated as follows:

$$z_t = \sigma(W_z x_t + U_z h_{t-1} + b_z) \quad (3)$$

The candidate hidden state (\tilde{h}_t) is calculated as

$$\tilde{h}_t = \tanh(W x_t + U(r_t \odot h_{t-1} + b_z)) \quad (4)$$

Finally, the output of the GRU (h_t) is calculated as a result of the previous output (h_{t-1}) and the candidate hidden state (\tilde{h}_t)

The graphical summary of the above described steps can be found in Figure 2

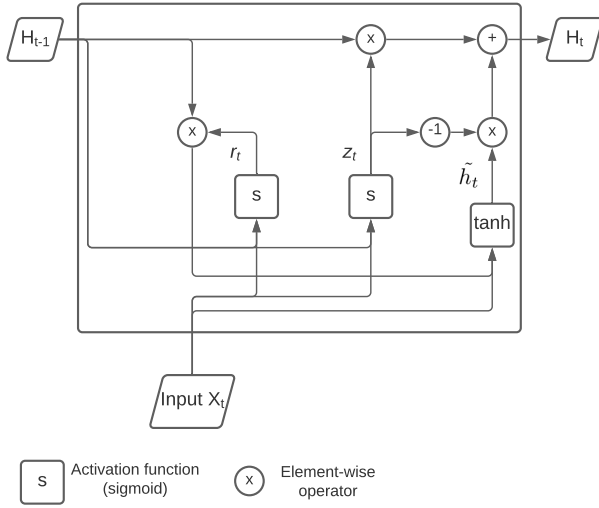


Figure 2. Gated Recurring Unit neural network architecture used in this study

To form a neural network from the GRU, the architecture represented in Figure 3 was implemented. In that image it is possible to appreciate that the input data (Open, Low, High and Close prices, along with the traded volume) as

a sequence is used to feed the neural network, this data is processed by multiple GRU units in the hidden layer, the output of these units is then directed to a fully connected that, produces five outputs (Open, Low, High, Close and Volume)

$$h_t = (1 - z_t)h_{t-1} + z_t\tilde{h}_t \quad (5)$$

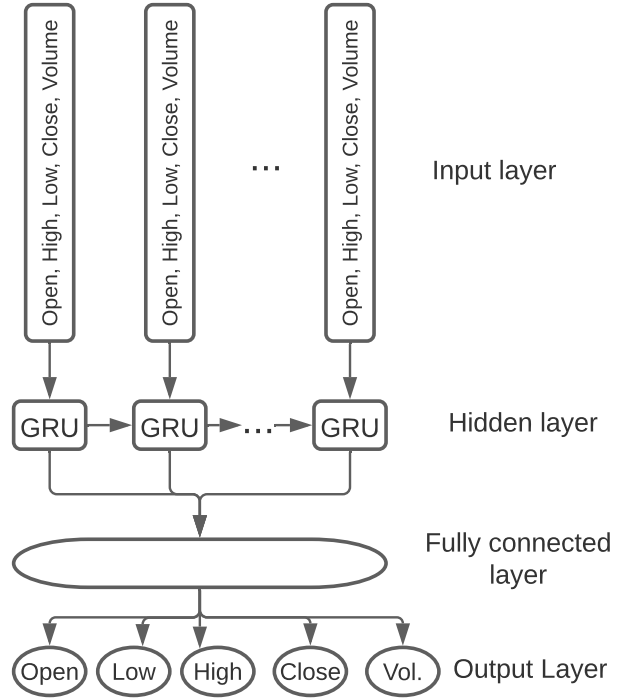


Figure 3. Gated Recurring Unit neural network architecture used in this study

3.4. Loss function and optimization criterion

In order to train the neural network, a Mean Square Error (MSE) loss function was used:

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (6)$$

where y is the predicted stock price and \hat{y} the true data

Using this loss function, backpropagation was used to update the parameters of the neural network, for the this purpose, the Adam optimizer was applied. The Adam optimizer was introduced by Kingma and Ba [6] and provides a computationally efficient method for stochastic optimization suitable for noise and sparse gradients. This optimizer is often used in the training of recurrent neural networks [11]

3.5. Hyperparameter tuning

In order to increase the prediction performance of the neural networks, several hyperparameters were tuned. To perform this, a validation set was created using the 20 last records of the training set (the same length as the test dataset). Using grid search (testing all the possible combination of parameters), the prediction performance of the neural network was tested on this validation dataset, and the MSE was calculated. The set of parameters with the smaller MSE was then selected for the final model.

In Table 1 it is possible to see that one of the hyperparameters (window size) is belong to the preprocessing of the data in this particular problem, while the other parameters belong to the architecture of the neural network (hidden layer size, number of recurrent layers) and to the training process (number of epochs and learning rate)

Hyperparameter	Values
Window size	20, 30, 40
Hidden Layer Size	20, 30, 40
Number of recurrent layers	1, 2, 3
Number of epochs	500, 1000, 1500
Learning rate	0.01, 0.05, 0.001

Table 1. Parameter tuning

3.6. Next day and intra day prediction

In order to test different input/output features and to increase the usability of the predictions of the neural network, three slightly different neural networks were tested. Firstly, the five features of the dataset were used as input and output (as in Figure 3), and the values of these features were predicted on the next day after the last observation on the input data (nextday model). Secondly, the same architecture was tested removing the traded Volume from the input and output data, as Wang *et al* [10] point out that the inclusion of this variable does not help in the prediction of the Stock Price (nextday w/o Volume).

Finally, as the trading operations are performed when the particular stock market is open (between Open and Close), a third variant, that uses the Open price to predict the Low, High and Close price on the same day was implemented. In this case, the input data is the Low, High and Close price of the N days before, along with the Open price of the same day and of the $N - 1$ days before. This last variant was developed to increase the usability of the models, as it can improve the trading decision making information during the day, as having the Low, High and Close prices can help the trader decide to buy or sell the stock before the end of the day.

4. Code and processing

The code, along with the requirements and setup instructions for the reproduction of the above-mentioned method, can be found in the following GitHub repository: <https://github.com/juliancabezas/RNN-GoogleStockPrice>. All the processing was performed in a Linux OS with NVIDIA GeForce GTX 1050 with CUDA support.

5. Results and discussion

5.1. Data correction

The correction of the Close price resulted in the prices shown in Figure 4, where is it possible to see that in the long term, the Google stock prices tends to increase its price, but showing several price drops in between. Lookign at the data, it is also possible to appreciate that the four prices are relatively close to each other, not showing a large intraday variations.




Figure 4. Open, High, Low and Close prices of the Google stock (corrected)

On the other hand, as seen in Figure 5, the traing volume presents an erratical and noisy behaviour, with several spikes in trading.

5.2. Hyperparameter tuning

The hyperparameters of the three models were (nextday, nextday w/o volume and intraday), showing the results of Table 2




Figure 5. Trading volume of the Google stock

Hyperparameter	Nextday	Nextday w/o vol.	Intraday
Window size	40	30	40
Hidden Layer Size	30	30	40
Number of recurrent layers	1	1	2
Number of epochs	1500	1500	1500
Learning rate	0.05	0.05	0.01

Table 2. Parameter tuning results for the three tested models

5.3. Training and testing of the final models

6. Conclusion

7. Bonus

This research provided insight over the importance of data augmentation m

References

- [1] Ayodele A. Adebisi, Aderemi O. Adewumi, and Charles K. Ayo. Stock price prediction using the ARIMA model. *Proceedings - UKSim-AMSS 16th International Conference on Computer Modelling and Simulation, UKSim 2014*, (March):106–112, 2014.
- [2] Kyunghyun Cho, Bart van Merriënboer, Dzmitry Bahdanau, and Yoshua Bengio. On the Properties of Neural Machine Translation: EncoderDecoder Approaches. *CoRR*, abs/1409.1:103–111, 2014.
- [3] Junyoung Chung, Caglar Gulcehre, KyungHyun Cho, and Yoshua Bengio. Empirical Evaluation of Gated Recurrent Neural Networks on Sequence Modeling. *CoRR*, abs/1412.3:1–9, 2014.
- [4] Sepp Hochreiter and Jürgen Schmidhuber. Long Short-Term Memory. *Neural Computation*, 9(8):1735–1780, 1997.
- [5] Yakup Kara, Melek Acar Boyacıoglu, and Ömer Kaan Baykan. Predicting direction of stock price index movement using artificial neural networks and support vector machines: The sample of the Istanbul Stock Exchange. *Expert Systems with Applications*, 38(5):5311–5319, 2011.
- [6] Diederik P. Kingma and Jimmy Lei Ba. Adam: A method for stochastic optimization. *3rd International Conference on Learning Representations, ICLR 2015 - Conference Track Proceedings*, pages 1–15, 2015.
- [7] Jiayu Qiu, Bin Wang, and Changjun Zhou. Forecasting stock prices with long-short term memory neural network based on attention mechanism. *PLoS ONE*, 15(1):1–15, 2020.
- [8] Arjun Singh Saud and Subarna Shakya. Analysis of look back period for stock price prediction with RNN variants: A case study on banking sector of NEPSE. *Procedia Computer Science*, 167(2019):788–798, 2020.
- [9] Sandro Skansi. *Introduction to Deep Learning: From Logical Calculus to Artificial Intelligence*. Springer, Cham, Switzerland, 2018.
- [10] Xiaohua Wang, Paul Kang Hoh Phua, and Weidong Lin. Stock Market Prediction Using Neural Networks: Does Trading Volume Help in Short-term Prediction? *Proceedings of the International Joint Conference on Neural Networks*, 4:2438–2442, 2003.
- [11] Jinghua Zhao, Dalin Zeng, Shuang Liang, Huilin Kang, and Qinming Liu. Prediction model for stock price trend based on recurrent neural network. *Journal of Ambient Intelligence and Humanized Computing*, 2020.

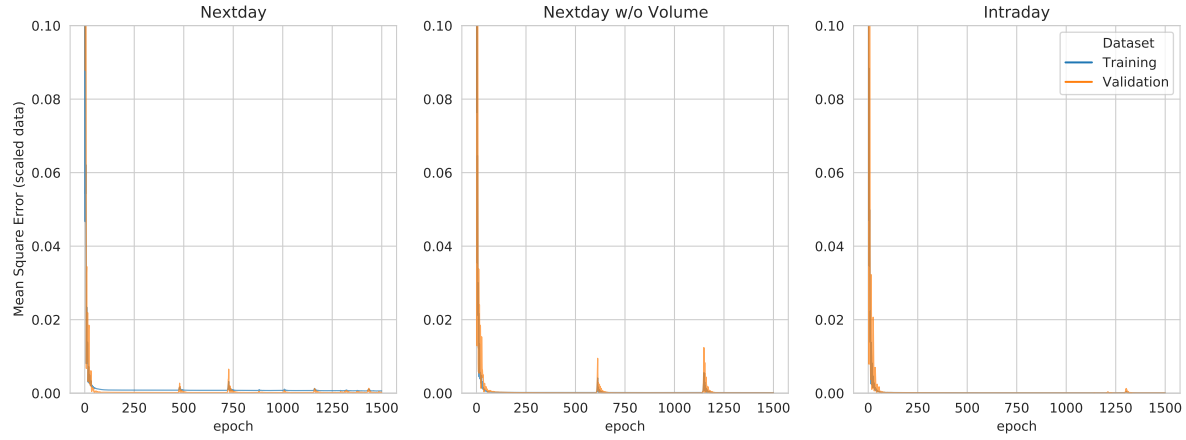


Figure 6. Training and validation loss curves for the three tested models

Variable	Nextday		Nextday w/o volume		Intraday	
	Train RMSE	Test RMSE	Train RMSE	Test RMSE	Train RMSE	Test RMSE
Open price	2.8498	5.9382	1.5054	5.5476	-	-
Low price	2.6912	6.8495	2.2645	6.6231	1.8879	3.9610
High price	3.7242	6.9893	2.6478	6.8147	1.9452	4.5980
Close price	2.9339	8.3147	2.9052	8.1642	2.4615	6.3192
Volume	954948.9845	1314561.2588	-	-	-	-
price						

Table 3. Performance of the models in the test dataset




Figure 7. Test data prediction of the trading volume of the Google stock (nextday model)




Figure 8. Prediction of the nextday model for the test data




Figure 9. Prediction of the nextday w/o volume for the test data

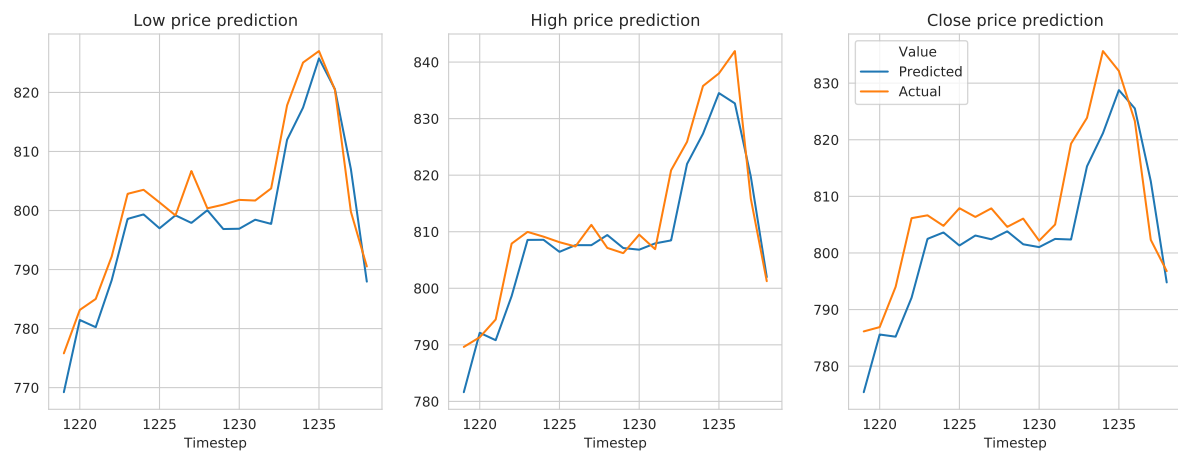


Figure 10. Prediction of the nextday w/o volume for the test data